

## Times Series Forecasting


We are supposed to use the Gas dataset inside the forecast package. We read the Gas data as a time series object. Now, it's time to do some exploratory analysis with the data.

We find that the data is a monthly data (frequency = 12).

We plot the month plot, season plot & the overall plot with the time series data.


### Month Plot

The month plot shows that the gas production has increased for every month of the year. We can also see a trend.




### Season Plot

The season plot shows a strong presence of seasonality.



## Time Series Plot

We have made a plot of the time series data. We can observe that till 1970, the graph is level and after that we can see a trend starting from 1970 to 1996. We can see a trend & seasonality.



We found that the periodicity is monthly.

## Test & Train Data

We divide the time series data into test & train data. Below is the code snippet for the same:

```
gas.train<- window(gas1, end = c(1994,12))
```

```
str(gas.train)
```

```
gas.test<- window(gas1, start = c(1995,1))
```

```
str(gas.test)
```


## Time Series Decomposition

We have done additive decomposition of the Time series in order to understand the components better:

```
gasComp = decompose(gas1) #additive
```

```
plot(gasComp)
```

## Decomposition of additive time series



We can see, that our gas time series data has:

- Level
- Trend
- Seasonality
- Random Component

### Holt-Winters' Model – Forecasting

We can see from decomposition that our time series data contains level, trend & seasonality. Hence we are applying Holt Winter's model for forecasting. It has three parameters: alpha, beta & gamma, where alpha signifies Trend, beta signifies Trend and gamma signifies seasonality.

```
HoltWinters.gas = HoltWinters(gas.train, alpha = NULL, beta = NULL, gamma = NULL,
```

```
seasonal = c("additive", "multiplicative"))
```

```
HW.forecast = forecast(HoltWinters.gas)
```

```
HW.forecast
```

```
HW.forecast$model
```

```
HW.forecast$mean
```

```
HW.forecast$fitted
```


```
accuracy(HW.forecast, gas.test)
```

```
autoplot(HW.forecast)
```

Below are the values of alpha, beta & gamma:

```
Smoothing parameters:  
alpha: 0.3766451  
beta : 0.009235467  
gamma: 0.7605161
```

Below is the graph showing forecasts from holt winter:




### Stationary Series Test for ARIMA Forecasting

Now, we will do ARIMA Forecasting. For applying ARIMA forecasting, we need a Stationary series. We have already established that the series has seasonality, but still we can check a stationary series by ADFTest.

We have done the ACF plot. From the ACF plot, we can determine that the series is seasonal.

### **Series gas1**



```
library(tseries)
```

```
autoplot(gas1)
```

```
acf(gas1, lag.max = 50)
```

```
adf.test(gas1)
```

We have to make Null Hypothesis and Alternate Hypothesis for ADF test.

Ho: Series not Stationary

H1: Series Stationary

```
> adf.test(gas1)
```

Augmented Dickey-Fuller Test

```
data: gas1
```

```
Dickey-Fuller = -2.7131, Lag order = 7, p-value = 0.2764
```

```
alternative hypothesis: stationary
```

Since the p value is 0.2764 we have to accept the null hypothesis, that is the series is non stationary.

Now, we will remove the seasonality of the time series by differencing of the series. We will do a differencing by 12 since seasonal series.

```
diff.TS=diff(diff(gas1,lag =12),1)
```

```
autoplot(diff.TS)
```

```
adf.test(diff.TS)
```


```
acf(diff.TS, lag.max=50)
```

```
acf(diff.TS, lag.max=50, plot = FALSE)
```

```
pacf(diff.TS, lag.max=50)
```

```
pacf(diff.TS, lag.max=50, plot = FALSE)
```

We get the following series by differencing.



Now, we do the ADF test again to check for stationarity.

```
> adf.test(diff.TS)
```

Augmented Dickey-Fuller Test


data: diff.TS

Dickey-Fuller = -9.0548, Lag order = 7, p-value = 0.01


alternative hypothesis: stationary

The low value of p value suggests that the series is stationary. The ACF test is done for the difference series. Now we can apply ARIMA Forecasting.

### Series diff.TS



### Series diff.TS



By the ACF and the PACF plot, we determine the value of  $p, d, q$  of  $ARIMA(p, d, q)$

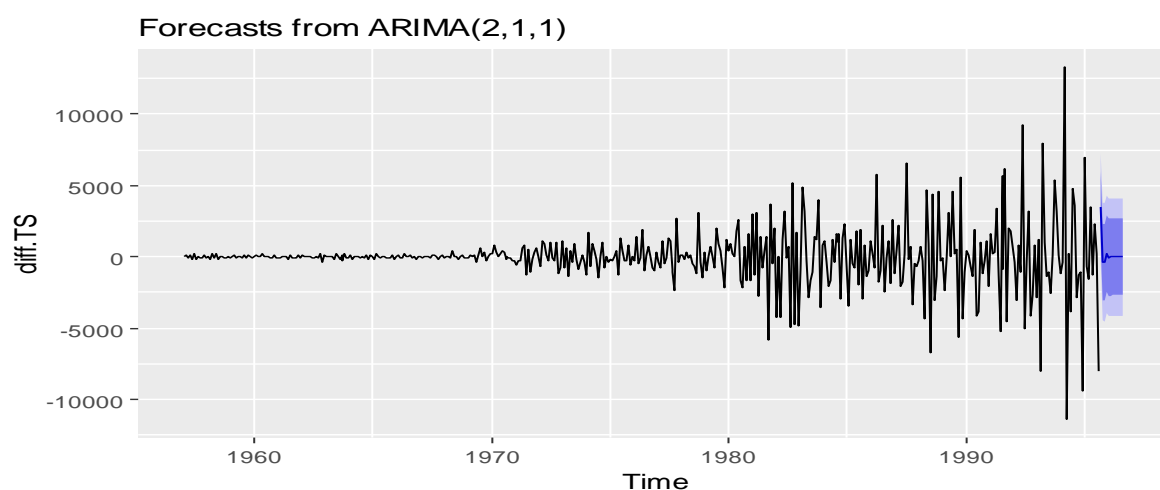
We will consider  $p = 2; d = 1; q = 1$  for the ARIMA Modeling.

### ARIMA Forecasting

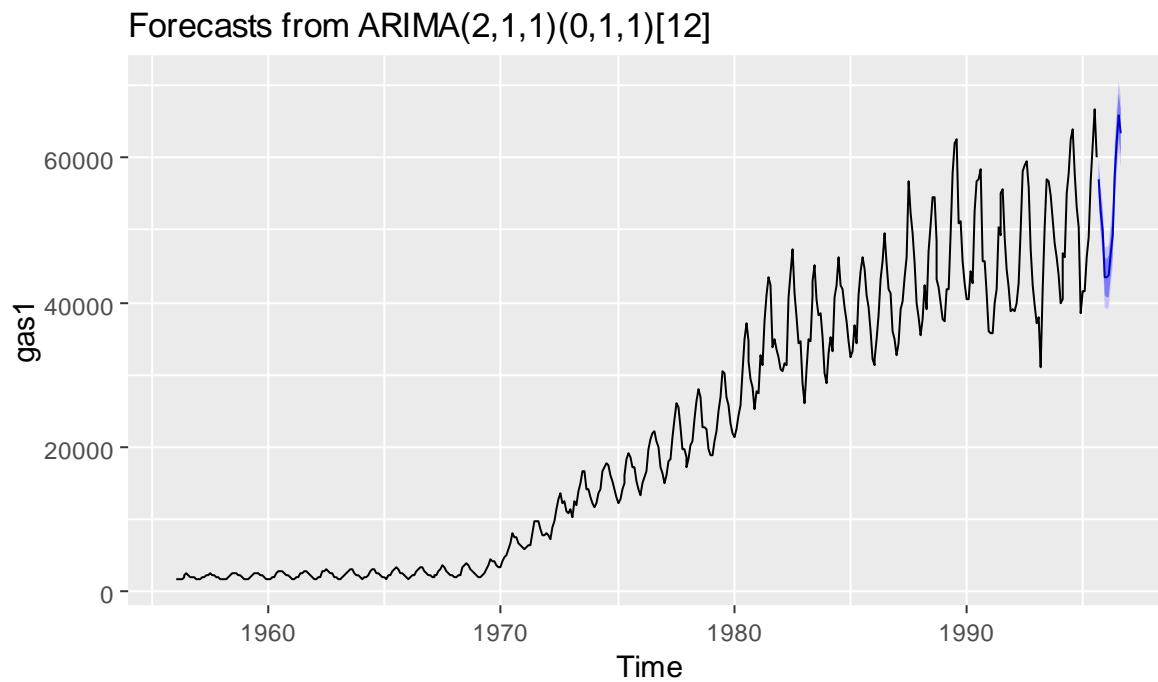
Now, since we have established that the diff.TS series is stationary.

```
library(fpp2)
gas.arima.fit <- Arima(diff.TS, order=c(2,1,1))
gas.arima.fit
checkresiduals(gas.arima.fit)
Box.test(gas.arima.fit$residuals, lag=30, type="Ljung-Box")
autoplot(forecast(gas.arima.fit, h=12))
gas.auto.arima.fit <- auto.arima(gas.train)
ARIMA.forecast=forecast(gas.auto.arima.fit)
accuracy(ARIMA.forecast, gas.test)
gas.auto.arima.fit <- auto.arima()
gas.auto.arima.fit <- auto.arima(gas1)
autoplot(forecast(gas.auto.arima.fit, h=12))
```

We will forecast for the next 12 periods using ARIMA Forecasting.



We have used the `auto.arima()` function to forecast the model. The below graph shows the forecasting.



### ACCURACY of the model

We have tested the `auto.arima` model with the test & train data set. We have found the Mean Absolute Percentage Error(MAPE) of the test set = 4.5%

```
> accuracy(ARIMA.forecast,gas.test)
```

	ME	RMSE	MAE	MPE	MAPE	MASE
ACF1 Theil's U						
Training set	18.43311	1551.733	876.3621	0.1845762	3.887880	0.4723883
.007149862	NA					
Test set	1797.05962	2623.714	2420.3779	3.4661214	4.504051	1.3046641
.417516251	0.5264559					